

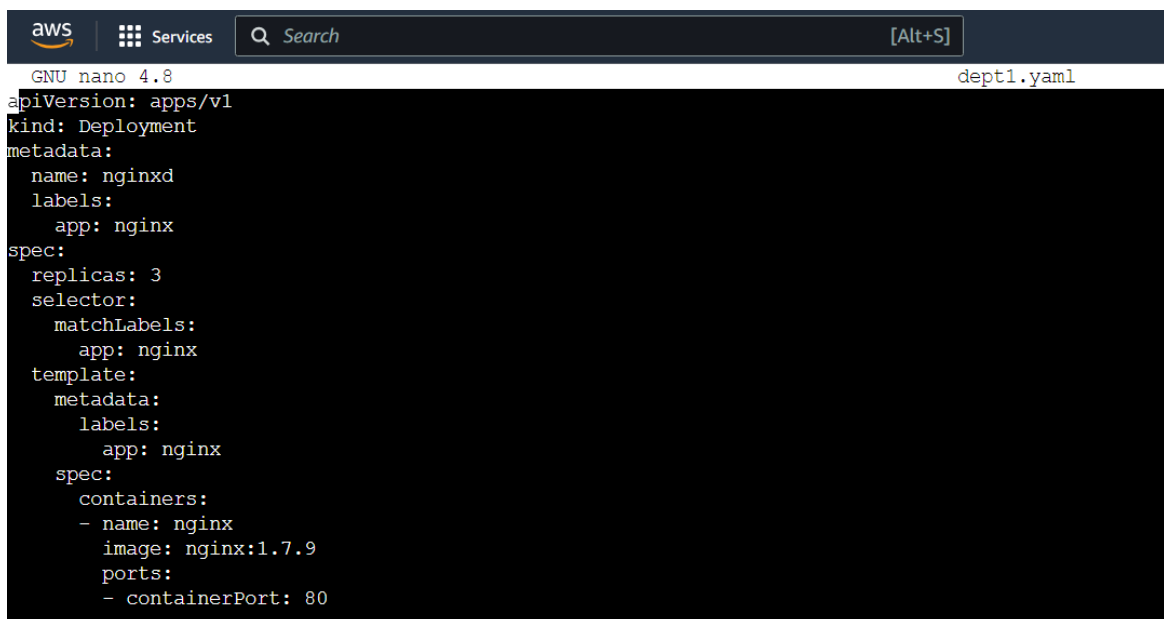
# Creating a Deployment

**Operation 1:** Write a yaml file for the Deployment that you are creating

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginxd
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
        ports:
        - containerPort: 80
```

Go to command line and create a yaml file and paste the above created specs.

**nano <filename>.yaml**



```
aws  Services  Search  [Alt+S]
GNU nano 4.8  dept1.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginxd
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.7.9
        ports:
        - containerPort: 80
```

Once done hit Ctrl+s and then Ctrl+x to save & exit

**Operation 2:** next thing to do is create the deployment yaml file.

`kubectl create -f <filename>`

```
ubuntu@ip-172-31-90-123:~$ nano dept1.yaml
ubuntu@ip-172-31-90-123:~$ kubectl create -f dept1.yaml
deployment.apps/nginxd created
ubuntu@ip-172-31-90-123:~$
```

Check if pods are created or not

`kubectl get pods`

```
ubuntu@ip-172-31-90-123:~$ nano dept1.yaml
ubuntu@ip-172-31-90-123:~$ kubectl create -f dept1.yaml
deployment.apps/nginxd created
ubuntu@ip-172-31-90-123:~$ kubectl get pods
NAME                                READY   STATUS            RESTARTS   AGE
nginx-deployment-86dcfdf4c6-vqldx  1/1     Terminating     0           45h
nginxd-9d6cbcc65-gpz76              1/1     Running           0           3m1s
nginxd-9d6cbcc65-nlrbl              1/1     Running           0           3m1s
nginxd-9d6cbcc65-svnps              1/1     Running           0           3m1s
ubuntu@ip-172-31-90-123:~$
```

**Operation 3:** You can have a look at the deployment, as well as the the replica set and the pods the deployment looks after

`kubectl get deploy`

```
ubuntu@ip-172-31-90-123:~$ kubectl get deploy
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
nginxd    3/3     3            3           6m15s
ubuntu@ip-172-31-90-123:~$
```

Or

`kubectl get deployment`

```
ubuntu@ip-172-31-90-123:~$ kubectl get deployment
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
nginxd    3/3     3            3           6m53s
ubuntu@ip-172-31-90-123:~$
```

**Operation 4:** now if I want to check replica sets

kubectl get rs

```
ubuntu@ip-172-31-90-123:~$ kubectl get rs
NAME                                DESIRED    CURRENT    READY    AGE
nginxd-9d6cbcc65                    3          3          3        9m5s
ubuntu@ip-172-31-90-123:~$
```

**Operation 5:** now scale your deployments using the command below

kubectl scale --replicas=<no of replica you want> deployments <deployment name>

```
ubuntu@ip-172-31-90-123:~$ kubectl scale --replicas=4 deployments nginxd
deployment.apps/nginxd scaled
ubuntu@ip-172-31-90-123:~$
```

Now let's check our pods

kubectl get pods

```
ubuntu@ip-172-31-90-123:~$ kubectl get pods
NAME                                READY    STATUS      RESTARTS    AGE
nginx-deployment-86dcfdf4c6-vqldx   1/1      Terminating    0           45h
nginxd-9d6cbcc65-9htpk               1/1      Running         0           114s
nginxd-9d6cbcc65-gpz76               1/1      Running         0           16m
nginxd-9d6cbcc65-nlrbl               1/1      Running         0           16m
nginxd-9d6cbcc65-svnps               1/1      Running         0           16m
ubuntu@ip-172-31-90-123:~$
```

**Operation 6:** change the image of the deployment your deployments using the command below

kubectl set image depolyments/<deployment name> <container name> = <image name>

```
ubuntu@ip-172-31-90-123:~$ kubectl set image deployments/nginxd nginx=centos:07
deployment.apps/nginxd image updated
ubuntu@ip-172-31-90-123:~$
```

To check the image use `kubectl describe deployment` command

```
Labels:  app=nginx
Containers:
  nginx:
    Image:      centos:07
    Port:       80/TCP
    Host Port:  0/TCP
    Environment: <none>
    Mounts:      <none>
    Volumes:      <none>
```

**Operation 7:** Now will check how rolling update works

`kubectl rollout status deployments/<deployment name>`

```
ubuntu@ip-172-31-90-123:~$ kubectl rollout status deployments/nginx
Waiting for deployment "nginx" rollout to finish: 2 out of 4 new replicas have been updated...
```

Command to check the status and the history the rollout that has been done –

`kubectl rollout history deployment/<deployment name>`

```
ubuntu@ip-172-31-90-123:~$ kubectl rollout history deployment/nginx
deployment.apps/nginx
REVISION  CHANGE-CAUSE
1          <none>
2          <none>

ubuntu@ip-172-31-90-123:~$
```

1) First we changed the no. replica

2) Second, we have changed the image name

**Operation 8:** now let's suppose the image centos:07 is not working properly now, and we want to retrieve back our older image, for that just write the below command-

`kubectl rollout undo deployment/<deployment name> --to-revision=<Revision no>`

```
ubuntu@ip-172-31-90-123:~$ kubectl rollout undo deployment/nginx --to-revision=1
deployment.apps/nginx rolled back
ubuntu@ip-172-31-90-123:~$
```

To check-

```
kubectl describe deployment <deployment name>
```

```
Labels:  app=nginx
Containers:
  nginx:
    Image:          nginx:1.7.9
    Port:           80/TCP
    Host Port:      0/TCP
    Environment:    <none>
    Mounts:         <none>
Volumes:          <none>
```

And our rollout status is also done

```
ubuntu@ip-172-31-90-123:~$ kubectl rollout status deployments/nginxd
deployment "nginxd" successfully rolled out
ubuntu@ip-172-31-90-123:~$
```