

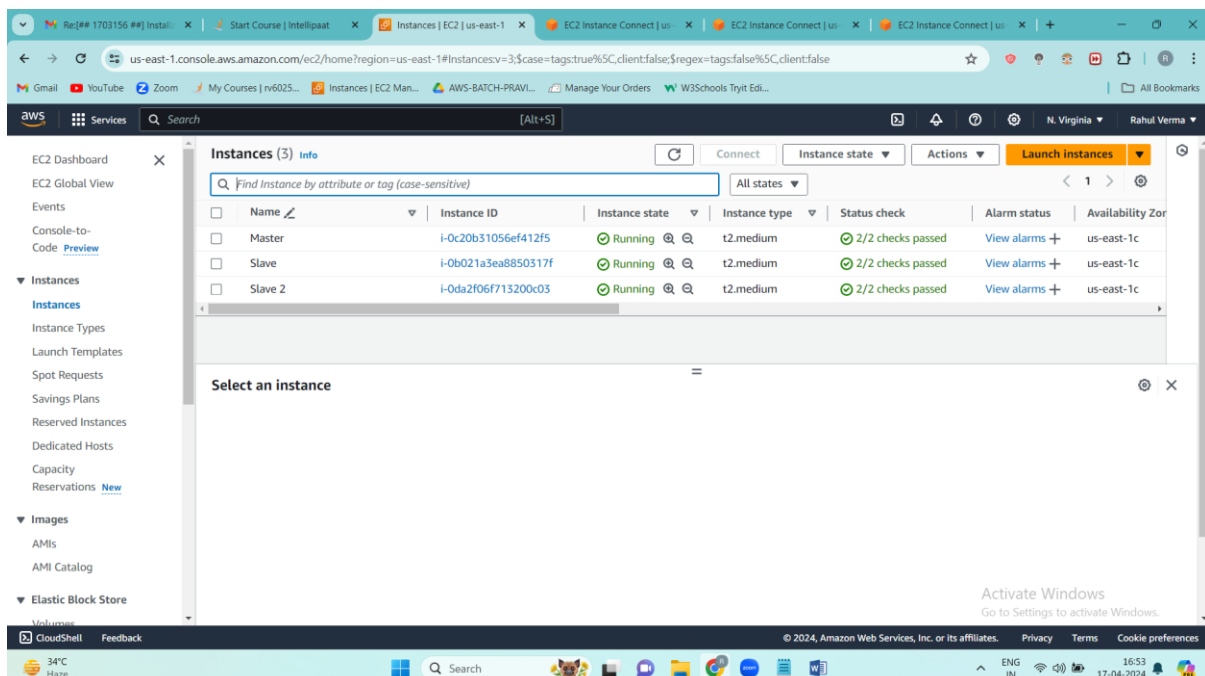
Hands-on

Installation of Kubernetes

Prerequisite:

Create 3 ec2 instances on any region with t2.medium instance type with Ubuntu OS.

1. Master
2. Slave
3. Slave 2



- Make sure all traffic is allowed on both the instance

Step 1: connect to your instances, Execute on all the instances i.e "Master" , "Slave" and "Slave 2"

Write-

```
sudo apt update
```

```
sudo nano install.sh
```

 (this will open nano editor)

now copy the below commands-

```
sudo apt-get update
```

```
sudo apt install docker.io -y
```

```
sudo apt-get install -y apt-transport-https ca-certificates curl  
gpg
```

```
sudo mkdir -p -m 755 /etc/apt/keyrings
```

```
curl -fsSL
```

```
https://pkgs.k8s.io/core:/stable:/v1.28/deb/Release.key |  
sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-  
keyring.gpg
```

```
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-  
keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.28/deb/ /' |  
sudo tee /etc/apt/sources.list.d/kubernetes.list
```

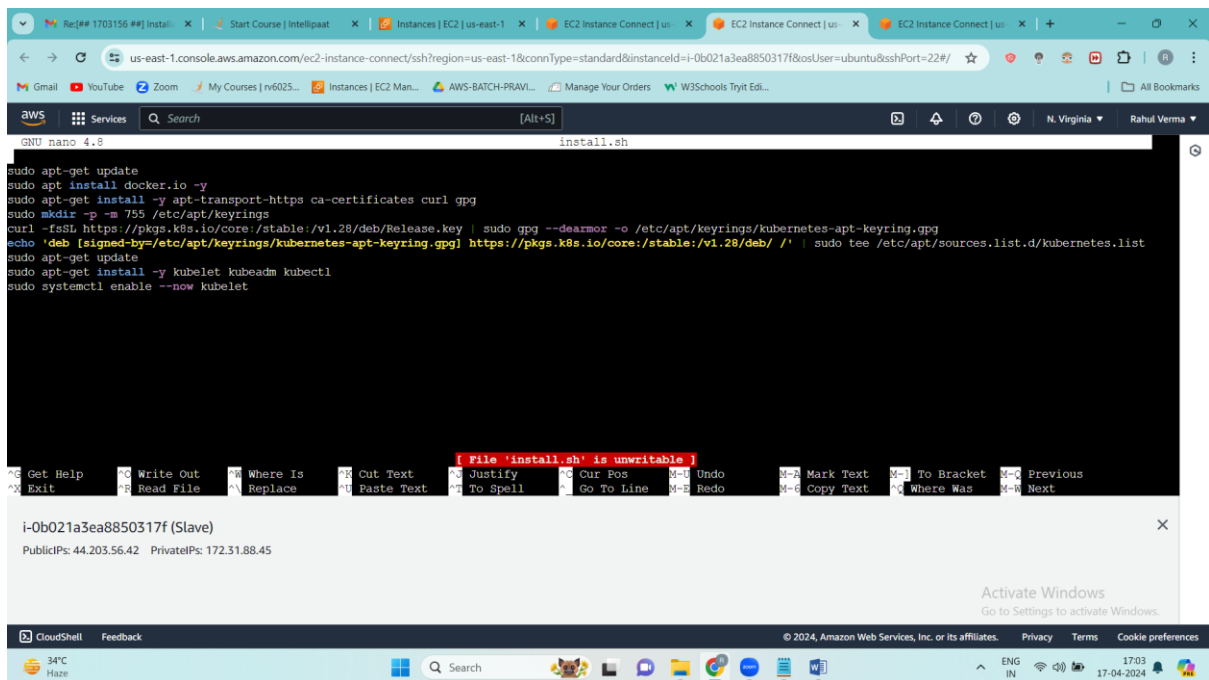
```
sudo apt-get update
```

```
sudo apt-get install -y kubelet kubeadm kubectl
```

```
sudo systemctl enable --now kubelet
```

Paste it in nano editor. Save and exit from it

Press ctrl+s and ctrl+x



The screenshot shows a terminal window within an AWS EC2 instance. The terminal is running the nano text editor to create a file named 'install.sh'. The commands being executed are:

```
sudo apt-get update
sudo apt install docker.io -y
sudo apt-get install -y apt-transport-https ca-certificates curl gpg
sudo mkdir -p -m 755 /etc/apt/keyrings
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.28/deb/Release.key | sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.28/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list
sudo apt-get update
sudo apt-get install -y kubelet kubeadm kubectl
sudo systemctl enable --now kubelet
```

The terminal output shows the successful execution of these commands. The nano editor's status bar at the bottom indicates that the file 'install.sh' is unwritable. The terminal window is titled 'i-0b021a3ea8850317f (Slave)' and shows the public and private IP addresses.

Write the below command, it execute all the commands from our install.sh file

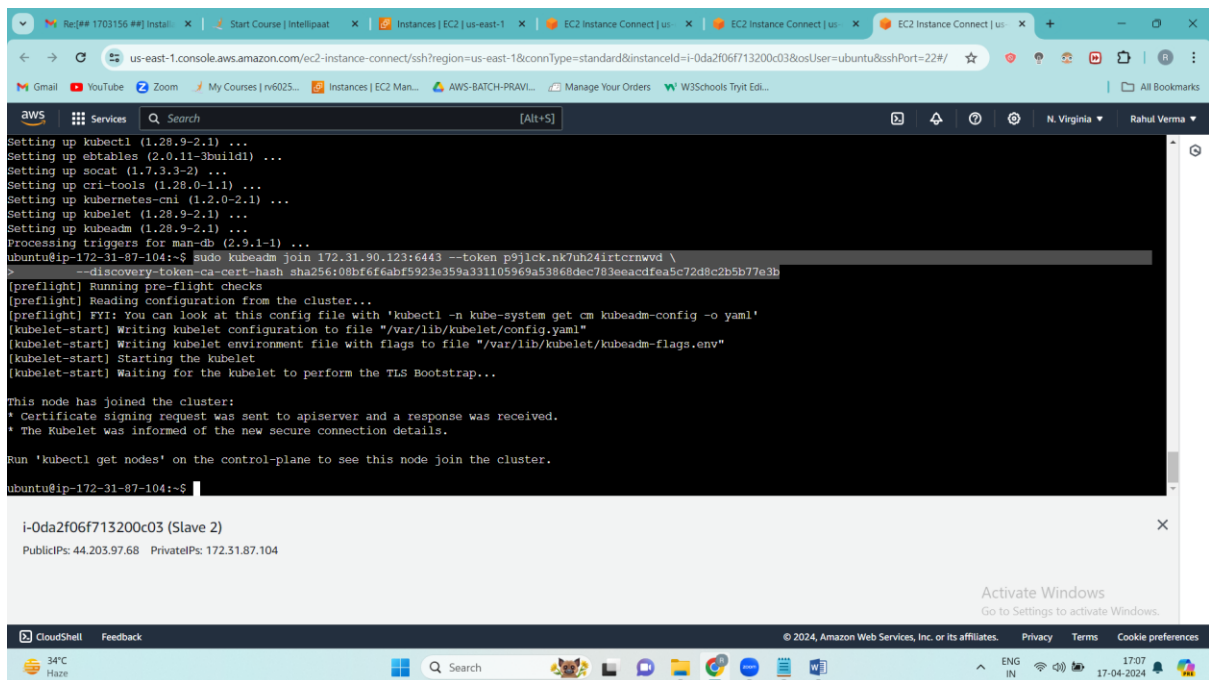
bash install.sh

Step 2: Execute this command on master NODE only

sudo kubeadm init --apiserver-advertise-address=privateipofmaster

Note: You need to replace “private_ip_of_master” with the actual private ip of your kubernetes master.

Paste the Token on Slave and Slave2



The screenshot shows a terminal window in AWS CloudShell. The terminal output indicates that various Kubernetes components (kubectl, ebttables, socat, cri-tools, kubelet, kubeadm) are being installed. A highlighted command is used to join the node to the cluster using a token and discovery token CA certificate hash. The output shows the node successfully joining the cluster, with preflight checks passing and the kubelet starting. The terminal also displays the instance details for 'i-Oda2f06f713200c03 (Slave 2)'.

```
Setting up kubectl (1.28.9-2.1) ...
Setting up ebttables (2.0.11-3build1) ...
Setting up socat (1.7.3.3-2) ...
Setting up cri-tools (1.28.0-1.1) ...
Setting up kubeadm (1.28.9-2.1) ...
Setting up kubelet (1.28.9-2.1) ...
Processing triggers for man-db (2.9.1-1) ...
ubuntu@ip-172-31-87-104:~$ sudo kubeadm join 172.31.90.123:6443 --token p9jick.nk7uh24irtcrnwvd \
--discovery-token-ca-cert-hash sha256:08bf6feabf5923e359a331105969a53868dec783eeacdf5c72d8c2b5b77e3b
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

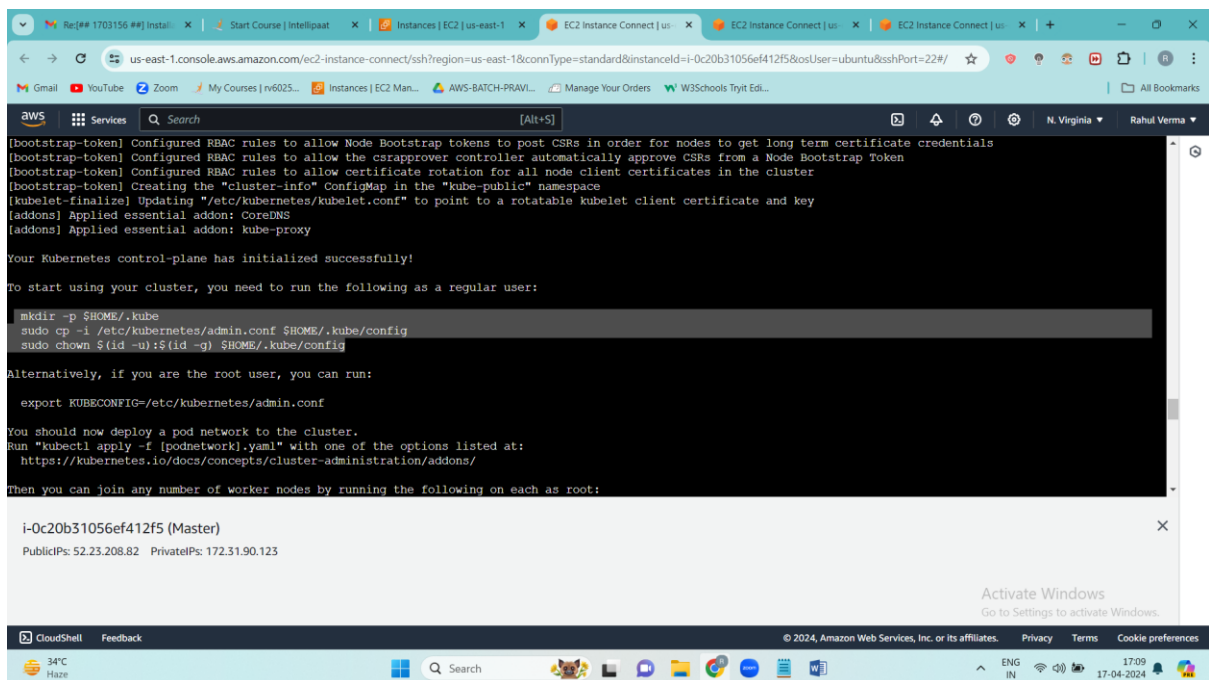
This node has joined the cluster:
* Certificate signing request was sent to apiserer and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
ubuntu@ip-172-31-87-104:~$
```

i-Oda2f06f713200c03 (Slave 2)
PublicIPs: 44.203.97.68 PrivateIPs: 172.31.87.104

Step3: Execute on Master node only

Copy the highlighted commands and paste it in master node one by one



The screenshot shows a terminal window in AWS CloudShell. The terminal output indicates that the Kubernetes control plane is being initialized. A highlighted command is used to create the kubeconfig file. The output shows the control plane initialized successfully, and the terminal also displays the instance details for 'i-Oc20b31056ef412f5 (Master)'.

```
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap Tokens to post CSRs in order for nodes to get long term certificate credentials
[bootstrap-token] Configured RBAC rules to allow the csrapprover controller automatically approve CSRs from a Node Bootstrap Token
[bootstrap-token] Configured RBAC rules to allow certificate rotation for all node client certificates in the cluster
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace
[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certificate and key
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run 'kubectl apply -f [podnetwork].yaml' with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:
```

i-Oc20b31056ef412f5 (Master)
PublicIPs: 52.23.208.82 PrivateIPs: 172.31.90.123

To start using your cluster, you need to run the following as a regular user:

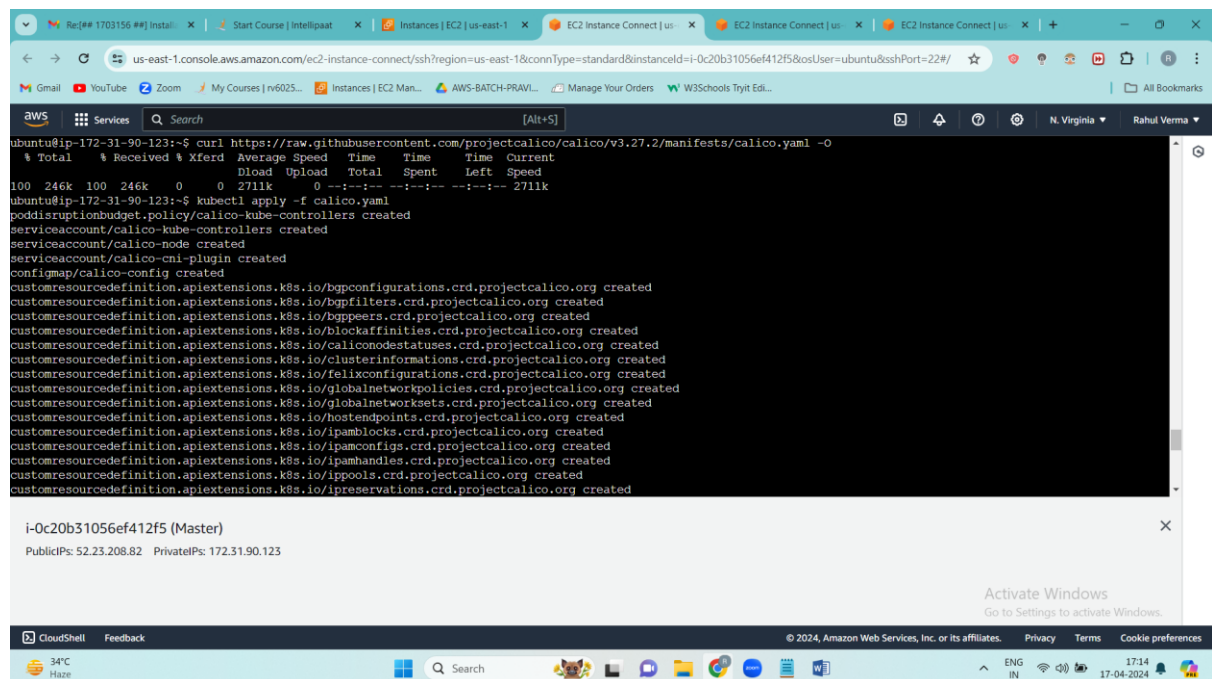
```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Step 4: Execute on master node only - Installing Calico

curl

<https://raw.githubusercontent.com/projectcalico/calico/v3.27.2/manifests/calico.yaml> -O

kubectl apply -f calico.yaml



The screenshot shows a terminal window on an AWS EC2 instance. The terminal output displays the following commands and their results:

```
ubuntu@ip-172-31-90-123:~$ curl https://raw.githubusercontent.com/projectcalico/calico/v3.27.2/manifests/calico.yaml -O
% Total % Received % Xferd Average Speed Time Time Time Current
Dload Upload Total Spent Left Speed
100 246k 100 246k 0 0 2711k 0 --:--:-- --:--:-- --:--:-- 2711k
ubuntu@ip-172-31-90-123:~$ kubectl apply -f calico.yaml
serviceaccount/calico-kube-controllers created
serviceaccount/calico-node created
serviceaccount/calico-cni-plugin created
configmap/calico-config created
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgpfilters.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipreservations.crd.projectcalico.org created
```

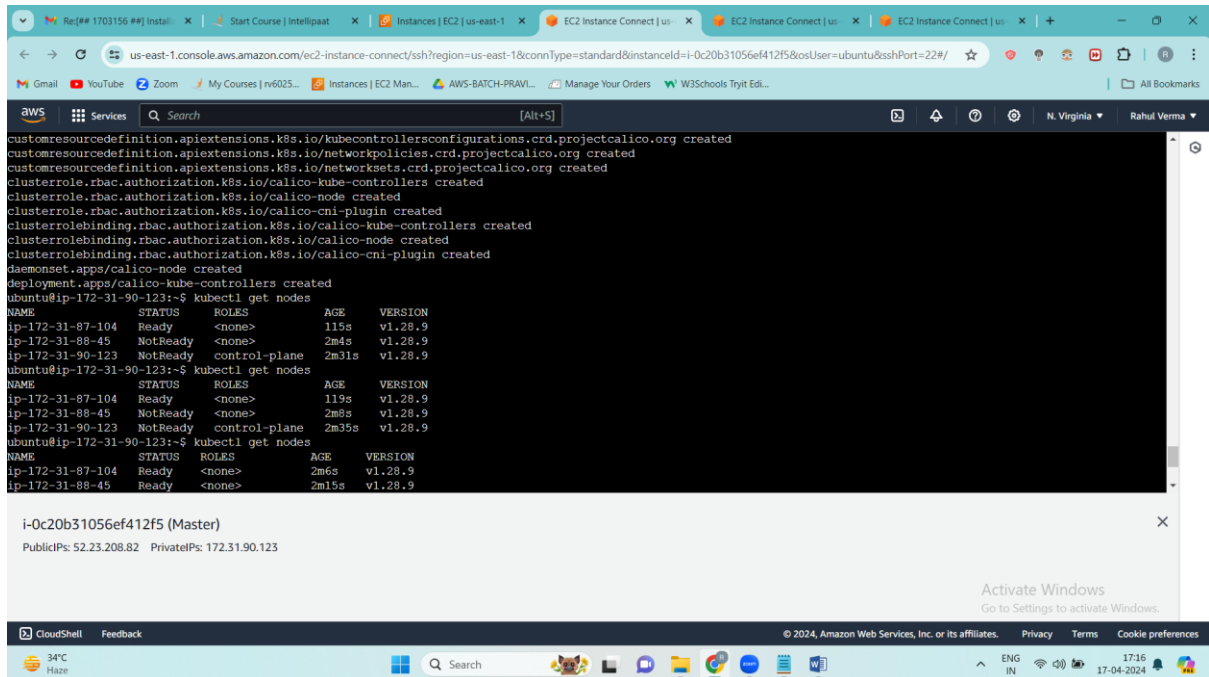
Below the terminal output, a box displays the instance details:

```
i-0c20b31056ef412f5 (Master)
PublicIPs: 52.23.208.82 PrivateIPs: 172.31.90.123
```

The bottom of the screenshot shows the Windows taskbar with the date and time as 17:14 on 17-04-2024.

Now to verify write the below command-

kubectl get nodes



The screenshot shows a terminal window in AWS CloudShell. The user has executed several kubectl commands to create resources and then run 'kubectl get nodes' twice. The output shows three nodes: ip-172-31-87-104 (Ready), ip-172-31-88-45 (NotReady), and ip-172-31-90-123 (NotReady, control-plane). The terminal also shows the creation of various Kubernetes resources like CustomResourceDefinition, ClusterRole, ClusterRoleBinding, and DaemonSet.

```
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-node created
clusterrole.rbac.authorization.k8s.io/calico-cni-plugin created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-cni-plugin created
daemonset.apps/calico-node created
deployment.apps/calico-kube-controllers created
ubuntu@ip-172-31-90-123:~$ kubectl get nodes
NAME                                STATUS    ROLES    AGE     VERSION
ip-172-31-87-104                    Ready     <none>    115s    v1.28.9
ip-172-31-88-45                    NotReady <none>    2m4s    v1.28.9
ip-172-31-90-123                    NotReady control-plane 2m31s    v1.28.9
ubuntu@ip-172-31-90-123:~$ kubectl get nodes
NAME                                STATUS    ROLES    AGE     VERSION
ip-172-31-87-104                    Ready     <none>    119s    v1.28.9
ip-172-31-88-45                    NotReady <none>    2m8s     v1.28.9
ip-172-31-90-123                    NotReady control-plane 2m35s    v1.28.9
ubuntu@ip-172-31-90-123:~$ kubectl get nodes
NAME                                STATUS    ROLES    AGE     VERSION
ip-172-31-87-104                    Ready     <none>    2m6s    v1.28.9
ip-172-31-88-45                    Ready     <none>    2m15s    v1.28.9
```

i-0c20b31056ef412f5 (Master)
PublicIPs: 52.23.208.82 PrivateIPs: 172.31.90.123

Activate Windows
Go to Settings to activate Windows.

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

34°C Haze ENG IN 17:16 17-04-2024