

# Creating a Taint on a node then a pod with tolerations but not to the taint

**Operation 1:** get the node name by running the command

`kubectl get nodes`

```
ubuntu@ip-172-31-90-123:~$ kubectl get nodes
NAME                                STATUS    ROLES    AGE   VERSION
ip-172-31-87-104                    Ready     <none>    2d19h v1.28.9
ip-172-31-88-45                     Ready     <none>    2d19h v1.28.9
ip-172-31-90-123                    Ready     control-plane 2d19h v1.28.9
ubuntu@ip-172-31-90-123:~$
```

Let me show you that we don't have any taints

`kubectl describe node <node name>`

```
projectcalico.org/IPv4IPPool
volumes.kubernetes.io/controller-managed-attach-detach
CreationTimestamp: Wed, 17 Apr 2024 11:14:43 +0000
Taints:            <none>
Unschedulable:    false
Lease:
  HolderIdentity:  ip-172-31-88-45
```

**Operation 2:** add a taint to a node using the following command

`kubectl taint nodes <node name> key=value:NoSchedule`

```
ubuntu@ip-172-31-90-123:~$ kubectl taint node ip-172-31-88-45 key=value:NoSchedule
node/ip-172-31-88-45 tainted
ubuntu@ip-172-31-90-123:~$
```

Now let's verify it

kubectl describe node <node name>

```
aws | Services | Search [Alt+S]
ubuntu@ip-172-31-90-123:~$ kubectl taint node ip-172-31-88-45 key=value:NoSchedule
node/ip-172-31-88-45 tainted
ubuntu@ip-172-31-90-123:~$ kubectl describe node ip-172-31-88-45
Name: ip-172-31-88-45
Roles: <none>
Labels: beta.kubernetes.io/arch=amd64
        beta.kubernetes.io/os=linux
        kubernetes.io/arch=amd64
        kubernetes.io/hostname=ip-172-31-88-45
        kubernetes.io/os=linux
Annotations: kubeadm.alpha.kubernetes.io/cri-socket: unix:///var/run/containerd
              node.alpha.kubernetes.io/ttl: 0
              projectcalico.org/IPv4Address: 172.31.88.45/20
              projectcalico.org/IPv4IPIPTunnelAddr: 192.168.41.128
              volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Wed, 17 Apr 2024 11:14:43 +0000
Taints: key=value:NoSchedule
Unschedulable: false
Lease:
```

So what taints does not allow is scheduling of new pods

**Operation 3:** next thing to do is to create the yaml file.

First will create a pod which does not have tolerations

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
  env: test
spec:
  containers:
  - name: nginx
    image: nginx
  imagePullPolicy: IfNotPresent
```

To create pod use the below command

```
kubectl create -f <file_name>.yaml
```

```
ubuntu@ip-172-31-90-123:~$ nano notaint.yaml
ubuntu@ip-172-31-90-123:~$ kubectl create -f notaint.yaml
pod/nginx created
ubuntu@ip-172-31-90-123:~$
```

Now if we will check the pod status it should be in pending status

```
kubectl get pods
```

```
ubuntu@ip-172-31-90-123:~$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
countdown-4hlgk                    0/1     Completed 0           17h
nginx                               0/1     Pending   0           2m54s
nginxd-9d6cbcc65-2zkvs             1/1     Running   1 (60m ago) 19h
nginxd-9d6cbcc65-6dwdv             1/1     Running   1 (60m ago) 19h
nginxd-9d6cbcc65-kfhd9             1/1     Running   1 (60m ago) 19h
rcsise-9xz8v                       1/1     Running   1 (60m ago) 19h
sharevol                           2/2     Running   2 (60m ago) 17h
ubuntu@ip-172-31-90-123:~$
```

So our taint is working properly

The screenshot shows the AWS Management Console interface. The main content area displays the details of a pod named 'nginx'. The pod is in a 'Pending' state. The 'Conditions' section shows 'PodScheduled' as 'False'. The 'Volumes' section shows 'Kube-api-access-98m49' as 'Projected'. The 'Events' section shows a 'Warning' of type 'FailedScheduling' with a message: '0/3 nodes are available: 1 node(s) had untoleration taint (node-role.kubernetes.io/control-plane: 1); 2 node(s) had untoleration taint (key: value), preemption: 0/3 nodes are available: 3 Preemption is not helpful for scheduling..'. The pod is scheduled on the master node 'i-0c20b31056ef412f5'.

```
Port: <none>
Host Port: <none>
Environment: <none>
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-98m49 (ro)
Conditions:
  Type             Status
  PodScheduled     False
Volumes:
  kube-api-access-98m49:
    Type:              Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:      kube-root-ca.crt
    ConfigMapOptional:  <nil>
    DownwardAPI:        true
QoS Class:           BestEffort
Node-Selectors:       <none>
Tolerations:          node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                      node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type    Reason             Age   From          Message
  ----    -
  Warning  FailedScheduling  76s (x3 over 11m)  default-scheduler  0/3 nodes are available: 1 node(s) had untoleration taint (node-role.kubernetes.io/control-plane: 1); 2 node(s) had untoleration taint (key: value), preemption: 0/3 nodes are available: 3 Preemption is not helpful for scheduling..
ubuntu@ip-172-31-90-123:~$
```

i-0c20b31056ef412f5 (Master)  
PublicIPs: 54.163.144.253 PrivateIPs: 172.31.90.123

## Now let's try to create a pod that uses tolleration

**Operation 4:** For that let's create a yaml file

`nano <file name>.yaml`

copy the below code

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx1
  labels:
    env: test
spec:
  containers:
  - name: nginx1
    image: nginx
    imagePullPolicy: IfNotPresent
  tolerations:
  - key: "key"
    operator: "Equal"
    value: "value"
    effect: "NoSchedule"
```

And now create pod using that yaml file

`kubectl create -f <file name>.yaml`

```
ubuntu@ip-172-31-90-123:~$ nano taint.yaml
ubuntu@ip-172-31-90-123:~$ kubectl create -f taint.yaml
pod/nginx1 created
ubuntu@ip-172-31-90-123:~$
```

Now this pod should be scheduled because it uses toleration

Let's check

```
kubectl get pods
```

```
ubuntu@ip-172-31-90-123:~$ nano taint.yaml
ubuntu@ip-172-31-90-123:~$ kubectl create -f taint.yaml
pod/nginx1 created
ubuntu@ip-172-31-90-123:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
countdown-4hlgk	0/1	Completed	0	18h
nginx	0/1	Pending	0	23m
nginx1	1/1	Running	0	11s
nginxd-9d6cbcc65-2zkvs	1/1	Running	1 (81m ago)	19h
nginxd-9d6cbcc65-6dwdv	1/1	Running	1 (81m ago)	19h
nginxd-9d6cbcc65-kfhd9	1/1	Running	1 (81m ago)	19h
rcsise-9xz8v	1/1	Running	1 (81m ago)	19h
sharevol	2/2	Running	2 (81m ago)	17h

```
ubuntu@ip-172-31-90-123:~$
```

And if we do

```
kubectl describe nginx1
```

The screenshot shows a terminal window with the output of the command `kubectl describe nginx1`. The output includes details about the pod's status, conditions, volumes, node selectors, tolerations, and recent events.

```
Started: Sat, 20 Apr 2024 07:47:08 +0000
Ready: True
Restart Count: 0
Environment: <none>
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-c4b9d (ro)
Conditions:
  Type             Status
  Initialized       True
  Ready             True
  ContainersReady   True
  PodScheduled      True
Volumes:
  kube-api-access-c4b9d:
    Type: Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI: true
QoS Class: BestEffort
Node-Selectors: <none>
Tolerations:
  key=value:NoSchedule
  node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
  node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type    Reason      Age    From          Message
  ----    -
  Normal  Scheduled   9m45s  default-scheduler  Successfully assigned default/nginx1 to ip-172-31-87-104
  Normal  Pulled      9m44s  kubelet        Container image "nginx" already present on machine
  Normal  Created     9m44s  kubelet        Created container nginx1
  Normal  Started     9m44s  kubelet        Started container nginx1
ubuntu@ip-172-31-90-123:~$
```